

Yatin Kande

Dearborn, MI (Open to Relocate) | +1 313-413-8327 | yatinrags135@gmail.com | linkedin/yatin-kande | github/YatinKande | yatinkande-portfolio.vercel.app

Summary

AI/ML Engineer with an M.S. in Data Science, focused on production systems across computer vision, generative AI, and predictive modeling. Recent work includes a Retrieval-Augmented Generation (RAG) system that answered diagnostic queries in under 8 seconds at 84% retrieval precision, a churn model that identified at-risk customers 30 days before they churned while achieving 81% AUC-ROC, and an Apache Spark pipeline that decreased processing time by 95%.

Education

University of Michigan-Dearborn

M.S. in Data Science | GPA: 3.77

Jain University

B.S. (Honors) in Data Science

Dearborn, MI

Aug 2024 – Dec 2025

Bengaluru, India

Jul 2020 – May 2023

Technical Skills

Languages & Libraries: Python, SQL, R, Pandas, NumPy, Scikit-learn, XGBoost, PyTorch, TensorFlow, Matplotlib, Seaborn

ML & Statistics: Regression, Classification, Clustering, Hypothesis Testing, Feature Engineering, Hyperparameter Tuning, SMOTE

Deep Learning & GenAI: CNN, 3D CNN, YOLO, ResNet, BERT, RAG, LangChain, FAISS, Hugging Face, Prompt Engineering, LLM Fine-Tuning

Cloud & Tooling: AWS (S3, Lambda, SageMaker, DynamoDB, SQS, Lex), Azure, GCP, Apache Spark, Microsoft Copilot Studio, Power Platform, Docker, MLflow, Git, CI/CD

Professional Experience

Software Engineer Intern, AI/ML (Part-Time)

Jan 2026 – Present

Fortisys, Inc.

Pleasanton, CA (Remote)

- Automated engagement, comment, and sentiment tracking across Instagram, Facebook, and X, with LinkedIn integration planned, by building AI agents in Copilot Studio and Power Platform for the social media team.
- Gave stakeholders weekly, per-platform visibility into engagement and sentiment by building a SharePoint reporting pipeline, later extending it with a Power Automate flow and Power BI dashboard for reads, impressions, and comments.
- Expanding automated outreach beyond social media by building, with the team, a new autonomous agent in Copilot Studio and Power Platform to recruit prospective clients and providers for the company.
- Applying the same Copilot Studio and Power Platform foundation used for the social media agents to this new client- and provider-recruiting use case, currently in active development.

ML Data Analyst Intern

Feb 2024 – Jul 2024

DataZymes

Bengaluru, India

- Reached 74% mean Average Precision (mAP) across 12 crop disease classes, up from a 51% single-stage baseline, by building a YOLOv5 pipeline trained on 10,000+ field images.
- Cut batch-processing runtime 95%, from over 4 hours to 18 minutes, by rebuilding the pipeline with Apache Spark and OpenCV.
- Sped up computer vision inference 40% by replacing the single-frame loop with batched processing and edge-level optimizations.
- Lifted the minority-class F1 score from 0.72 to 0.76 by growing the training set with 1,500+ synthetic images from augmentation and web scraping.
- Cut weekly review prep from 3 hours to 20 minutes by building Power BI dashboards used by 8 engineers across 5 model iterations.

Machine Learning Engineer Intern

Mar 2022 – Aug 2022

SmartKnower

Bengaluru, India

- Gave the retention team churn risk scores they could act on within 24 hours by building an XGBoost model on 8,500+ records that scored 81% AUC-ROC.
- Gained 4% precision over Random Forest and Logistic Regression by tuning hyperparameters with GridSearchCV across learning rate, depth, and regularization.
- Lifted baseline accuracy 5% by engineering 15+ behavioral features from raw temporal data.
- Cut the false-negative rate on high-risk churners from 18% to 12% by applying SMOTE to boost minority-class recall 6%.

Projects

AutoRAG: Automotive Diagnostic AI

Nov 2025 – Jan 2026

- Reached 84% retrieval precision, well above the 61% dense-only baseline, by building a hybrid RAG system over service manuals and OBD-II images.
- Improved answer quality 12% over single-retriever setups by combining FAISS and BM25 through LangChain with Gemini 2.0 Flash.
- Cut diagnostic query time from 45 seconds to under 8 seconds by tuning chunk size, embedding strategy, and the reranking pipeline.
- Kept diagnostic support available without internet access by adding Ollama Llama 3.2 Vision as an offline fallback.

Dataset Discovery Chatbot

Aug 2025 – Oct 2025

- Cut dataset search time from 5-10 minutes to under 10 seconds by launching a serverless chatbot on AWS Lex and Lambda.
- Reduced repeat searches for 50+ users by designing a DynamoDB-backed scoring engine that ranks matches against session history.
- Replaced manual polling with real-time email alerts by automating dataset delivery through an S3, SQS, and SNS event pipeline.
- Cut overall dataset search time 70% across 50+ queries spanning 6 topic domains.

Certifications

- OCI 2025 Certified AI Foundations Associate (Oracle, 2025)
- LangChain for LLM Application Development (DeepLearning.AI, 2025)